

Imubit Challenge:

- **Dataset:** I created a new dataset contains 5K pairs of images from CIFAR10 with their appropriate pair of labels.

Few examples of the new dataset:



{'airplane': 0, 'automobile': 0, 'bird': 0, 'cat': 0, 'deer': 0, 'dog': 0, 'frog': 1, 'horse': 1, 'ship': 0, 'truck': 0}



{'airplane': 0, 'automobile': 0, 'bird': 1, 'cat': 1, 'deer': 0, 'dog': 0, 'frog': 0, 'horse': 0, 'ship': 0, 'truck': 0}



{'airplane': 0, 'automobile': 0, 'bird': 1, 'cat': 0, 'deer': 0, 'dog': 0, 'frog': 0, 'horse': 0, 'ship': 0, 'truck': 1}



- For all the following model, I split the train set to train and validation (80\20) and only check the final result (after hyperparameter-tuning) on the test set. I treated the test set as unknown throughout all the training.
I used data augmentation technique while training to prevent overfitting and to improve the generalization of the model.
- **Network Architecture:** I created a CNN model for the new dataset. The model basically contains 3 building blocks of:

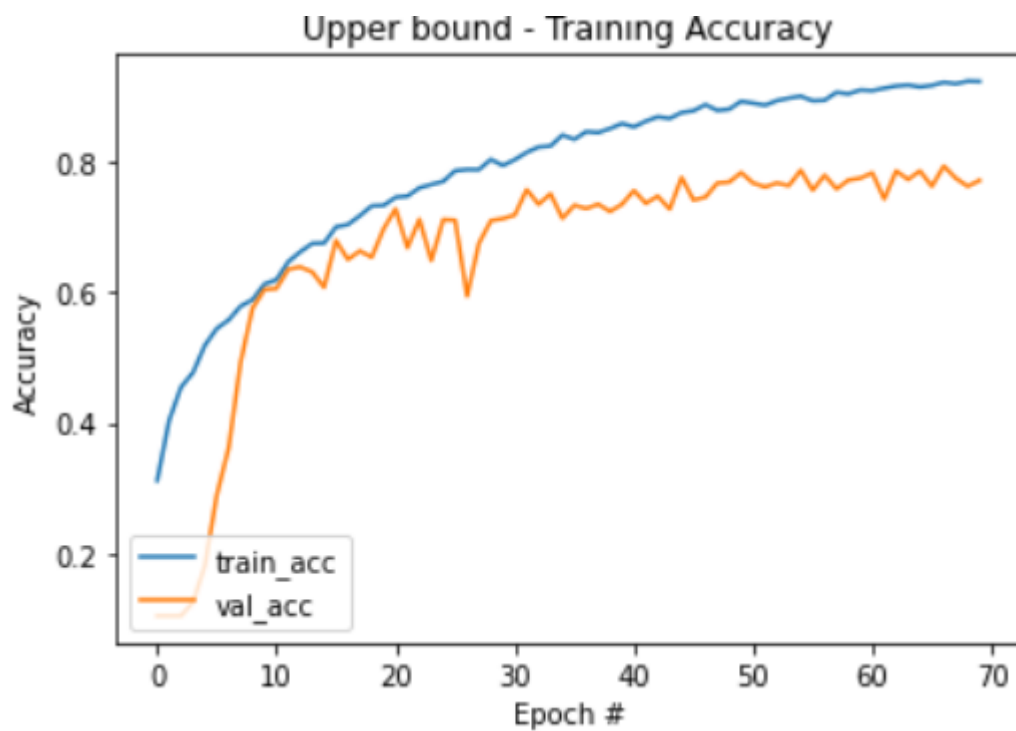
CONV2D -> BN -> CONV2D -> BN -> MaxPooling

On top of it I stick two fully-connected layer. I end the model with 10 units and a softmax activation, which is suitable for the multi-class classification.

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130
Total params: 1,343,018		


- I used an early stopping (according the validation loss) to decide when to stop the train.

- **Upper Bound Model:** I trained a CNN model on the total 10K images (5K * 2) and the **real** labels as upper bound to the expected accuracy in the challenge problem.



The loss and the accuracy graphs indicate a standard learning and convergence process of the network. The final accuracy (on the test set) is: 0.743



- **Naïve Baseline:** Next, I trained a naïve baseline. This network has same architecture as the upper bound network. It was trained on the pair-based data as follows, randomly assign one of the labels to one of the images. In this way, it is clear that the data has about fifty percent of incorrect labels.




The final accuracy (on the test set) is: 0.473, significantly lower than the accuracy of the upper bound. The network very fast overfitted the train set.

- **Solution:** In order to solve the challenge problem, I used a Siamese network, for the pair of images, with share weights. The loss is calculated twice- for the two possible permutations and the minimal option is chosen.

Schematic description:



The final accuracy (on the test set) is: 0.711, close to the upper bound. (If I take the validation from the real label, I reach to 0.740, very close to the upper bound)



I try also to initial my model with the weights from the baseline model, but it didn't improve the final results.

To conclude, my model outperforms the baseline model ($0.71 > 0.473$) and very close to the upper bound (0.743).

*There are more techniques to improve the overall performance like use more complicated network backbone or transfer learning, but I think it is out of the scope of this challenge.